



Table of Contents

- 1 Numba
 - 1.1 jit 与 njit
 - 1.2 Loops
 - 1.3 FastMath
 - 1.4 Parallel
- 2 JAX
 - 2.1 替换 NumPy
 - 2.2 jit
 - 2.3 grad
 - 2.4 vmap
- 3 Cython
 - 3.1 annotate
 - 3.2 cfunc/cdef
 - 3.3 性能对比
- 4 CuPy
 - 4.1 cupy.ndarray
 - 4.2 Device
 - 4.3 Data Transfer
 - 4.4 Memory
- 5 Sparse
 - 5.1 创建
 - 5.2 转换
 - 5.3 计算
- 6 Dask
 - 6.1 创建
 - 6.2 索引
 - 6.3 计算
 - 6.4 任务图
 - 6.5 底层 API
 - 6.6 调度
 - 6.7 性能对比
- 7 Xarray
 - 7.1 创建
 - 7.2 索引
 - 7.3 属性
 - 7.4 计算
 - 7.5 GroupBy
 - 7.6 可视化
- 8 小结
- 9 参考

```
In [1]: import numpy as np
        np.__version__
```

```
Out[1]: '1.22.4'
```

文档阅读说明：

-  表示 Tip
-  表示注意事项

本部分内容主要介绍与NumPy相关的高性能、分布式数值计算用法和工具。它们的安装都比较简单，参考文档即可。我们这里侧重介绍一下每个工具是干什么的，有什么特点，我们什么时候需要使用它们。

Numba

文档：[Numba documentation — Numba 0.55.2+0.g2298ad618.dirty-py3.7-linux-x86_64.egg documentation](#)

Numba是适用于Python的即时编译器，最适合用在使用NumPy数组和函数，以及循环的代码中。最常用的使用方法是通过装饰器。当调用一个Numba的装饰器时，它会被编译为「即时」的机器代码以供执行，全部或部分代码随后可以以机器代码的速度运行。

概括来说，这几种情况适合使用Numba：

- 很多数学计算
- 使用了很多Numpy
- 有很多循环

它的原理是，通过读取装饰函数的Python字节码，并将其与函数输入的参数类型信息相结合，分析和优化代码后，使用LLVM编译器根据CPU定制生成函数的机器代码版本。之后的调用都会使用该编译后的版本。

```
In [2]: from numba import jit
```

```
In [14]: def func_normal(a):
        x = np.median(a)
        y = np.max(a)
        t = x / y;
        z = x * np.sqrt(1 + t * t)
        m = 0.0
        for i in range(a.shape[0]):
            m += np.tanh(a[i, i])
            m /= z
        return a + m
```

```
In [6]: a = np.arange(100, dtype=np.int32).reshape(10, 10)
```

```
In [15]: %timeit func_normal(a)
```

68.2 μ s \pm 2.32 μ s per loop (mean \pm std. dev. of 7 runs, 10000 loops each)

```
In [16]: @jit(nopython=True)
def func_numba(a):
    x = np.median(a)
    y = np.max(a)
    t = x / y;
    z = x * np.sqrt(1 + t * t)
    m = 0.0
    for i in range(a.shape[0]):
        m += np.tanh(a[i, i])
        m /= z
    return a + m
```

```
In [19]: prebuild = func_numba(a)
```

```
In [20]: %timeit func_numba(a)
```

1.68 μ s \pm 33.8 ns per loop (mean \pm std. dev. of 7 runs, 1000000 loops each)

可以很明显看出性能的提升。

jit与njit

Numba 有两种模式：

- `nopython` 模式：用 `@jit(nopython=True)` 或 `@njit` 装饰器装饰。这种模式下，函数将完全在编译模式下运行，不需要Python解释器参与。这也是 Numba 推荐的使用方式。
- `object` 模式：直接用 `@jit` 装饰时，如果 `nopython` 模式失败，则会使用 `object` 模式进行编译，此时一部分可「Numba」的代码会使用机器代码执行，剩下的则使用Python编译器执行。

```
In [3]: from numba import njit
import pandas as pd
```

```
In [4]: @njit
def jit_fail(x):
    df = pd.DataFrame(x)
    df += 1
    cov = df.cov()
    return cov
```

```
In [5]: x = {'a': [1, 2, 3], 'b': [20, 30, 40]}
```

```
In [6]: jit_fail(x)
```

```

-----
TypeError                                Traceback (most recent call last)
<ipython-input-6-1a477a3676a1> in <module>
----> 1 jit_fail(x)

/usr/local/lib/python3.8/site-packages/numba/core/dispatcher.py in _compile_for_args(self, *args, **kws)
    399             e.patch_message(msg)
    400
--> 401             error_rewrite(e, 'typing')
    402         except errors.UnsupportedError as e:
    403             # Something unsupported is present in the user code, add help
info

/usr/local/lib/python3.8/site-packages/numba/core/dispatcher.py in error_rewrite(e, issue_type)
    342             raise e
    343         else:
--> 344             reraise(type(e), e, None)
    345
    346         argtypes = []

/usr/local/lib/python3.8/site-packages/numba/core/utils.py in reraise(tp, value, tb)
    78         value = tp()
    79         if value.__traceback__ is not tb:
---> 80             raise value.with_traceback(tb)
    81         raise value
    82

TypeError: Failed in nopython mode pipeline (step: nopython frontend)
non-precise type pyobject
[1] During: typing of argument at <ipython-input-4-481c7b069f0d> (3)

File "<ipython-input-4-481c7b069f0d>", line 3:
def jit_fail(x):
    df = pd.DataFrame(x)
    ^

This error may have been caused by the following argument(s):
- argument 0: cannot determine Numba type of <class 'dict'>

```

```

In [7]: @jit
def jit_succ(x):
    df = pd.DataFrame(x)
    df += 1
    cov = df.cov()
    return cov

```

```

In [8]: # 会有警告
cov = jit_succ(x)

```

```

<ipython-input-7-cb5f8409f203>:1: NumbaWarning:
Compilation is falling back to object mode WITH looplifting enabled because Function "jit_succ" failed type inference due to: non-precise type pyobject
[1] During: typing of argument at <ipython-input-7-cb5f8409f203> (3)

File "<ipython-input-7-cb5f8409f203>", line 3:
def jit_succ(x):
    df = pd.DataFrame(x)
    ^

    @jit
/usr/local/lib/python3.8/site-packages/numba/core/object_mode_passes.py:177: NumbaWarning: Function "jit_succ" was compiled in object mode without forceobj=True.

File "<ipython-input-7-cb5f8409f203>", line 2:
@jit
def jit_succ(x):
^

    warnings.warn(errors.NumbaWarning(warn_msg,
/usr/local/lib/python3.8/site-packages/numba/core/object_mode_passes.py:187: NumbaDeprecationWarning:
Fall-back from the nopython compilation path to the object mode compilation path has been detected, this is deprecated behaviour.

For more information visit http://numba.pydata.org/numba-doc/latest/reference/deprecation.html#deprecation-of-object-mode-fall-back-behaviour-when-using-jit

File "<ipython-input-7-cb5f8409f203>", line 2:
@jit
def jit_succ(x):
^

    warnings.warn(errors.NumbaDeprecationWarning(msg,

```

```
In [13]: %timeit jit_succ(x)
```

559 µs ± 16 µs per loop (mean ± std. dev. of 7 runs, 1000 loops each)

```
In [10]: def func(x):
        df = pd.DataFrame(x)
        df += 1
        cov = df.cov()
        return cov
```

```
In [14]: %timeit func(x)
```

514 µs ± 29.8 µs per loop (mean ± std. dev. of 7 runs, 1000 loops each)

此时性能差不多，Numba 反而会慢一些，因为它还要判断是不是可以编译优化。

[官方文档](#)还有其他一些特性，不过我们主要关注和性能相关的几个。

Loops

Numba 对循环可以进行优化：

```
In [116... def get_primes(x):
    res = []
    for v in range(x+1):
        if v < 2:
            continue
        flag = True
        for i in range(2, int(np.sqrt(v)) + 1):
            if v % i == 0:
                flag = False
        if flag:
            res.append(v)
    return res
```

```
In [117... is_prime(10)
```

```
Out[117... [2, 3, 5, 7]
```

```
In [118... x = 100000
```

```
In [119... %timeit get_primes(x)
```

1.28 s ± 48 ms per loop (mean ± std. dev. of 7 runs, 1 loop each)

```
In [120... @njit
def jit_get_primes(x):
    res = []
    for v in range(x+1):
        if v < 2:
            continue
        flag = True
        for i in range(2, int(np.sqrt(v)) + 1):
            if v % i == 0:
                flag = False
        if flag:
            res.append(v)
    return res
```

```
In [121... prebuild = jit_get_primes(x)
```

```
In [122... %timeit jit_get_primes(x)
```

77.7 ms ± 2.16 ms per loop (mean ± std. dev. of 7 runs, 10 loops each)

可以看出效果还是很明显的，不过我们在使用时尽量做一下性能对比，以做到心中有数。

FastMath

在某些情况下，可以通过放松一些严格的（IEEE754）数值获得额外的性能提升。

IEEE 二进制浮点数算术标准（IEEE 754）是 20 世纪 80 年代以来最广泛使用的浮点数运算标准，为许多 CPU 与浮点运算器所采用。这个标准定义了表示浮点数的格式（包括负零 -0）与反常值（denormal number），一些特殊数值（（无穷（Inf）与非数值（NaN）），以及这些数值的“浮点数运算符”；它也指明了四种数值舍入规则和五种例外状况（包括例外发生的时机与处理方式）。——维基百科

以官方文档例子来说明：

```
In [205... @njit(fastmath=False)
def do_sum(A):
    acc = 0.
    # without fastmath, this loop must accumulate in strict order
    for x in A:
        acc += np.sqrt(x)
    return acc

@njit(fastmath=True)
def do_sum_fast(A):
    acc = 0.
    # with fastmath, the reduction can be vectorized as floating point
    # reassociation is permitted.
    for x in A:
        acc += np.sqrt(x)
    return acc
```

```
In [206... a = np.arange(40000)
```

```
In [207... prebuild1 = do_sum(a)
prebuild2 = do_sum_fast(a)
```

```
In [208... prebuild1, prebuild2
```

```
Out[208... (5333233.1256554425, 5333233.1256554425)
```

```
In [209... %timeit do_sum(a)
```

70 μ s \pm 678 ns per loop (mean \pm std. dev. of 7 runs, 10000 loops each)

```
In [210... %timeit do_sum_fast(a)
```

53.1 μ s \pm 1.58 μ s per loop (mean \pm std. dev. of 7 runs, 10000 loops each)

默认情况下，编译器在浮点优化方面收到严格限制（如重新关联浮点表达式），因为这样的优化可能导致结果改变。比如：

- $(10000001.0f * 10000001.0f) / 10000001.0f == 10000000.0f$
- $10000001.0f * (10000001.0f / 10000001.0f) == 10000001.0f$

第一个表达式括号里的会超出32位精度，会舍入。

更多关于这方面的知识可以参考：

- [SIMD vectorization](#)
- [Floating Point Optimization](#)

Parallel

```
In [199... from numba import prange
```

```
In [211... @njit(parallel=True)
def do_sum_parallel(A):
```

```

# each thread can accumulate its own partial sum,
# and then a cross
# thread reduction is performed to obtain the result to return
n = len(A)
acc = 0.
for i in prange(n):
    acc += np.sqrt(A[i])
return acc

@njit(parallel=True, fastmath=True)
def do_sum_parallel_fast(A):
    n = len(A)
    acc = 0.
    for i in prange(n):
        acc += np.sqrt(A[i])
    return acc

```

```

In [212... prebuild1 = do_sum_parallel(a)
prebuild2 = do_sum_parallel_fast(a)
prebuild1, prebuild2

```

```

Out[212... (5333233.125655441, 5333233.125655442)

```

```

In [213... %timeit do_sum_parallel(a)

```

```

108 µs ± 1.3 µs per loop (mean ± std. dev. of 7 runs, 10000 loops each)

```

```

In [214... %timeit do_sum_parallel_fast(a)

```

```

95.5 µs ± 3.72 µs per loop (mean ± std. dev. of 7 runs, 10000 loops each)

```

JAX

文档: [JAX Quickstart — JAX documentation](#)

JAX是运行在CPU、GPU和TPU上的NumPy，两者关系：

- JAX提供了一个方便的受NumPy启发的接口。
- 通过鸭子类型，JAX数组通常可以直接替代NumPy数组。
- 与NumPy数组不同，JAX数组是不可变的。

替换NumPy

```

In [215... import jax.numpy as jnp

```

```

In [216... jnp.arange(10)

```

```

WARNING:absl:No GPU/TPU found, falling back to CPU. (Set TF_CPP_MIN_LOG_LEVEL=0 and rerun for more info.)

```

```

Out[216... DeviceArray([0, 1, 2, 3, 4, 5, 6, 7, 8, 9], dtype=int32)

```

```

In [217... list(_)

```

```

Out[217... [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]

```



```
In [218... from jax import random
```

```
In [221... key = random.PRNGKey(42)
```

```
In [223... a = random.normal(key, (2, 3))  
a
```

```
Out[223... DeviceArray([[ 0.61226517,  1.1225883 ,  1.1373315 ],  
                [-0.81273264, -0.8904051 ,  0.12623137]], dtype=float32)
```

```
In [224... b = random.normal(key, (3, 2))  
b
```

```
Out[224... DeviceArray([[ 0.61226517,  1.1225882 ],  
                [ 1.1373315 , -0.8127326 ],  
                [-0.8904051 ,  0.12623137]], dtype=float32)
```

```
In [225... jnp.dot(a, b)
```

```
Out[225... DeviceArray([[ 0.63893783, -0.08147554],  
                [-1.6226908 , -0.1727684 ]], dtype=float32)
```

```
In [226... np.dot(a, b)
```

```
Out[226... array([[ 0.63893783, -0.08147555],  
                [-1.6226908 , -0.1727684 ]], dtype=float32)
```

jit

jit 主要用来加速。

```
In [227... from jax import jit
```

```
In [232... def func_normal(a):  
    x = jnp.median(a)  
    y = jnp.max(a)  
    t = x / y;  
    z = x * jnp.sqrt(1 + t * t)  
    m = 0.0  
    for i in range(a.shape[0]):  
        m += jnp.tanh(a[i, i])  
    m /= z  
    return a + m
```

```
In [233... a = np.arange(100, dtype=np.int32).reshape(10, 10)
```

```
In [235... pre = jit(func_normal)(a)
```

```
In [238... %timeit func_normal(a)
```

5.83 ms ± 57.9 µs per loop (mean ± std. dev. of 7 runs, 100 loops each)

```
In [236... %timeit jit(func_normal)(a)
```

42.5 µs ± 2.45 µs per loop (mean ± std. dev. of 7 runs, 10000 loops each)

grad

`grad` 用来计算导数。

```
In [240... from jax import grad
```

以Sigmoid函数为例：

$$f(x) = \frac{1}{1 + e^{-x}}$$

它的导数为： `f(x) * (1-f(x))` 。

```
In [245... def sigmoid(x):  
    return 1.0 / (1.0 + jnp.exp(-x))
```

```
In [253... def dsigmoid(x):  
    x = sigmoid(x)  
    return x * (1-x)
```

```
In [254... dersigmoid = grad(sigmoid)
```

```
In [255... dersigmoid(2.)
```

```
Out[255... DeviceArray(0.10499357, dtype=float32)
```

```
In [256... dsigmoid(2.)
```

```
Out[256... DeviceArray(0.10499363, dtype=float32)
```

vmap

`vmap` 用于自动向量化或批量化。以官方文档为例：

```
In [258... mat = random.normal(key, (150, 100))
```

```
In [259... batched_x = random.normal(key, (10, 100))
```

首先看简单循环版：

```
In [260... def apply_matrix(v):  
    return jnp.dot(mat, v)
```

```
In [261... def naive_batched(v_batched):  
    return jnp.stack([apply_matrix(v) for v in v_batched])
```

```
In [265... naive_batched(batched_x).shape
```

```
Out[265... (10, 150)
```

```
In [275... %timeit naive_batched(batched_x).block_until_ready()
```

3.78 ms ± 168 µs per loop (mean ± std. dev. of 7 runs, 100 loops each)

接下来是矩阵乘法（手动Batch）版：

```
In [276... def batched_apply_matrix(v_batched):  
            return jnp.dot(v_batched, mat.T)
```

```
In [277... %timeit batched_apply_matrix(batched_x).block_until_ready()
```

226 μ s \pm 24.7 μ s per loop (mean \pm std. dev. of 7 runs, 1000 loops each)

最后是 `vmap`：

```
In [278... from jax import vmap, jit
```

```
In [281... @jit  
def vmap_apply_matrix(v_batched):  
    return vmap(apply_matrix)(v_batched)
```

```
In [282... %timeit vmap_apply_matrix(batched_x).block_until_ready()
```

18.9 μ s \pm 545 ns per loop (mean \pm std. dev. of 7 runs, 100000 loops each)

这在没法使用矩阵乘法的时候非常有用。

最后值得一提的是，三个方法乃至 `jnp` 既可以单独使用，也可以联合起来使用。实际使用时可以根据自己的需要灵活组合。

这里我们简单介绍一下，更多内容可以进一步阅读文档。

Cython

[Welcome to Cython's Documentation — Cython 3.0.0a10 documentation](#)

Cython在本章中都比较特别，它是一种编程语言，使得编写C扩展像Python一样容易。它旨在成为Python的超集，赋予它高级、面向对象和动态编程。Cython代码会被翻译成优化的C/C++代码并编译为Python扩展模块。不仅使得程序执行与C语言紧密集成，同时保持Python的易开发性。

看个最简单的例子：

```
In [5]: # 加载扩展  
%load_ext Cython
```

```
In [47]: %%cython  
  
cdef int a = 0  
for i in range(10):  
    a += i  
print(a)
```

45

annotate

可以使用 `annotate` 选项查看代码分析：

In [48]: `%%cython --annotate`

```
cdef int a = 0
for i in range(10):
    a += i
print(a)
```

45

Out[48]: Generated by Cython 0.29.30

Yellow lines hint at Python interaction.

Click on a line that starts with a " + " to see the C code that Cython generated for it.

```
1:
+2: cdef int a = 0
+3: for i in range(10):
+4:     a += i
+5: print(a)
```

下面显示的是纯Python版，不过这种情况下需要类型标记。

In [49]: `%%cython --annotate`

```
a: cython.int = 0
for i in range(10):
    a += i
print(a)
```

45

Out[49]: Generated by Cython 0.29.30

Yellow lines hint at Python interaction.

Click on a line that starts with a " + " to see the C code that Cython generated for it.

```
1:
+2: a: cython.int = 0
+3: for i in range(10):
+4:     a += i
+5: print(a)
```

当然，即便是纯Python代码，也可以使用Cython先编译，获得性能提升。不过对于性能关键的代码，添加静态类型声明通常很有用。

In [50]: `%%cython --annotate`

```
a = 0
for i in range(10):
    a += i
print(a)
```

Out[50]: Generated by Cython 0.29.30

Yellow lines hint at Python interaction.

Click on a line that starts with a "+" to see the C code that Cython generated for it.

```
1:
+2: a = 0
+3: for i in range(10):
+4:     a += i
+5: print(a)
```

cfunc/cdef

Python函数调用可能很耗时——在Cython中可能是双重的，因为可能需要在Python对象之间进行转换才能调用。因此Cython提供了声明C样式函数的方法，Cython特定的 `cdef` 语句，以及 `@cfunc` 装饰器用以在Python语法中声明C样式函数。两种方法会生成相同的C代码。

In [2]: `import cython`

In [3]: `@cython.cfunc`
`@cython.exceptval(-2, check=True)`
`def f(x: cython.double) -> cython.double:`
 `return x ** 2 - x`

In [6]: `%%cython`
`cdef double f(double x) except? -2:`
 `return x ** 2 - x`

性能对比

接下来，我们用实际例子来对比性能。

In [24]: `# 标准Python版`
`def get_primes(num):`
 `res = [0] * 1000`
 `v = 2`
 `len_res = 0`
 `while len_res < num:`
 `flag = True`
 `for i in range(2, int(np.sqrt(v)) + 1):`
 `if v % i == 0:`
 `flag = False`
 `if flag:`
 `res[len_res] = v`
 `len_res += 1`
 `v += 1`
 `return res`

```
In [7]: %%cython

cdef extern from "math.h":
    double sqrt(double x)

def cython_get_primes(int num):
    cdef int i, n, v=2, len_res=0
    cdef flag
    cdef int res[1000]

    while len_res < num:
        flag = True
        n = int(sqrt(v)) + 1
        for i in range(2, n):
            if v % i == 0:
                flag = False
        if flag:
            res[len_res] = v
            len_res += 1
        v += 1
    return res
```

```
In [30]: %timeit ps1 = cython_get_primes(1000)
2.12 ms ± 157 µs per loop (mean ± std. dev. of 7 runs, 100 loops each)
```

```
In [31]: %timeit ps2 = get_primes(1000)
37.4 ms ± 471 µs per loop (mean ± std. dev. of 7 runs, 10 loops each)
```

```
In [32]: ps1 == ps2
```

```
Out[32]: True
```

我们再对比一下直接编译后的Python代码，在code目录下执行：

```
python3 python setup.py build_ext --inplace
```

```
In [37]: # 导入
%cd code

/Users/Yam/Yam/powerful-numpy/src/skilled/code
```

```
In [38]: import primes
```

```
In [44]: ps3 = primes.get_primes(1000)
```

```
In [45]: ps3 == ps2
```

```
Out[45]: True
```

```
In [46]: %timeit primes.get_primes(1000)
23.1 ms ± 1.14 ms per loop (mean ± std. dev. of 7 runs, 10 loops each)
```

CuPy

文档: [CuPy – NumPy & SciPy for GPU — CuPy 10.5.0 documentation](https://docs.cupy.dev/en/stable/)

Cupy是一个兼容NumPy/SciPy的数组库，用于使用Python进行GPU加速计算。CuPy充当在NVIDIA CUDA或AMD ROCm平台上运行现有NumPy/SciPy代码的替代品。其主要目标是为Python用户提供GPU加速能力，无需深入了解底层GPU技术。

⚠ 注意，本节内容需要cuda环境。

```
In [2]: import cupy as cp
        cp.__version__
```

```
Out[2]: '10.5.0'
```

cupy.ndarray

```
In [3]: x_gpu = cp.array([1,2,3])
```

```
In [4]: x_gpu
```

```
Out[4]: array([1, 2, 3])
```

```
In [5]: type(x_gpu)
```

```
Out[5]: cupy._core.core.ndarray
```

`cupy.ndarray` 与 `np.ndarray` 的主要区别在于CuPy会把数组分配在当前设备（某一张GPU卡）上。其他的API都是和NumPy几乎没有区别的。如果熟悉NumPy，等于熟悉了CuPy。

```
In [6]: rng = cp.random.default_rng(42)
```

```
In [7]: rng.integers(0, 10, (2,3))
```

```
Out[7]: array([[5, 4, 0],
               [7, 2, 3]])
```

Device

这是CuPy比较重要的一个概念——「当前设备」，这是默认的GPU设备，数组的分配、操作和计算都在它上面运行。

```
In [8]: x_gpu0 = cp.array([1,2,3])
```

```
In [9]: x_gpu0.device
```

```
Out[9]: <CUDA Device 0>
```

```
In [10]: with cp.cuda.Device(1):
         x_gpu1 = cp.array([1,2,3])
```

```
In [11]: x_gpu1.device
```

```
Out[11]: <CUDA Device 1>
```

注意，这里你得有两张（或以上）的卡才行。比如我们再来一张不存在的卡：

```
In [12]: with cp.cuda.Device(2):  
         x_gpu2 = cp.array([1,2,3])
```

```
-----  
CUDARuntimeError                                Traceback (most recent call last)  
Input In [12], in <cell line: 1>()  
----> 1 with cp.cuda.Device(2):  
      2     x_gpu2 = cp.array([1,2,3])  
  
File cupy/cuda/device.pyx:184, in cupy.cuda.device.Device.__enter__()  
  
File cupy_backends/cuda/api/runtime.pyx:365, in cupy_backends.cuda.api.runtime.setDevice()  
  
File cupy_backends/cuda/api/runtime.pyx:142, in cupy_backends.cuda.api.runtime.check_status()  
  
CUDARuntimeError: cudaErrorInvalidDevice: invalid device ordinal
```

Data Transfer

主要指GPU卡和host（挂载卡的主机）之间的传输。

```
In [16]: x_cpu = np.array([1,2,3])
```

```
In [18]: type(x_cpu)
```

```
Out[18]: numpy.ndarray
```

```
In [19]: # 移动到GPU上  
         x_gpu = cp.asarray(x_cpu)
```

```
In [21]: type(x_gpu)
```

```
Out[21]: cupy._core.core.ndarray
```

`cp.asarray` 也可以在GPU卡之间互相移动。

```
In [22]: with cp.cuda.Device(1):  
         x_gpu2 = cp.asarray(x_gpu)
```

```
In [24]: x_gpu2.device
```

```
Out[24]: <CUDA Device 1>
```

注意，`cp.asarray` 不会复制数据，如果需要复制，可以使用 `cp.array(arr, dtype, copy=True)`。它实际上等价于 `cp.array(a, dtype, copy=False)`。

`copy=True` 会返回一个新数组，否则会返回对象。

```
In [40]: arr = cp.array([1,2,3])  
         cp.asarray(arr) is arr
```


Out[40]: True

```
In [41]: # 从GPU到Host
x_cpu2 = cp.asnumpy(x_gpu2)
x_cpu2
```

Out[41]: array([1, 2, 3])

```
In [42]: type(x_cpu2)
```

Out[42]: numpy.ndarray

```
In [43]: # 或者使用`get`方法
x_gpu2.get()
```

Out[43]: array([1, 2, 3])

```
In [44]: type(_)
```

Out[44]: numpy.ndarray

`cp.asnumpy` 返回NumPy数组（在Host上），`cp.asarray` 返回一个CuPy数组（在当前卡上）。两个方法都可以接受任意的输入（cp或np的数组）。

Memory

在GPU编程中，内存管理是个比较重要的环节。CuPy使用内存池管理内存，包括两种：

- Device内存池（GPU Memory），分配GPU内存时使用
- Pinned内存池（非交换CPU Memory），CPU到GPU数据传输时使用

```
In [3]: mempool = cp.get_default_memory_pool()
pinpool = cp.get_default_pinned_memory_pool()
```

```
In [4]: # 400bytes CPU内存
a_cpu = np.arange(100, dtype=np.float32)
```

```
In [5]: a_cpu.nbytes
```

Out[5]: 400

```
In [6]: mempool.used_bytes()
```

Out[6]: 0

```
In [7]: mempool.total_bytes()
```

Out[7]: 0

```
In [8]: pinpool.n_free_blocks()
```

Out[8]: 0

从CPU到GPU，一旦传输完成，pinned memory会被释放。

注意，实际分配的大小可能会四舍五入到大于请求大小的值。

```
In [9]: a = cp.array(a_cpu)
```

```
In [10]: a.nbytes
```

```
Out[10]: 400
```

```
In [11]: mempool.used_bytes()
```

```
Out[11]: 512
```

```
In [12]: mempool.total_bytes()
```

```
Out[12]: 512
```

```
In [13]: pinpool.n_free_blocks()
```

```
Out[13]: 1
```

如果数组超出所在域，GPU内存会被释放。

```
In [14]: a = None
```

```
In [15]: mempool.used_bytes()
```

```
Out[15]: 0
```

```
In [16]: mempool.total_bytes()
```

```
Out[16]: 512
```

```
In [17]: pinpool.n_free_blocks()
```

```
Out[17]: 1
```

使用 `free_all_blocks` 来清理内存池。

```
In [22]: mempool.free_all_blocks()
```

```
In [23]: mempool.used_bytes()
```

```
Out[23]: 0
```

```
In [24]: mempool.total_bytes()
```

```
Out[24]: 0
```

```
In [26]: pinpool.free_all_blocks()
```

```
In [27]: pinpool.n_free_blocks()
```

```
Out[27]: 0
```

CUDA编程中 `threads` , `blocks` 和 `grids` 是三个重要的概念:

- `thread`: 一个thread是运行在单个GPU核上的一系列指令
- `block`: 多个threads在GPU上以block的抽象单元执行
- `grid`: block的block又被称为grid

也可以对GPU的内存进行硬限制:

```
export CUPY_GPU_MEMORY_LIMIT="1073741824"
```

```
# or
```

```
export CUPY_GPU_MEMORY_LIMIT="50%"
```

或使用内置的方法:

```
In [55]: mempool = cp.get_default_memory_pool()
```

```
In [56]: with cp.cuda.Device(0):  
         mempool.set_limit(size=1024**3)
```

```
In [57]: cp.get_default_memory_pool().get_limit()
```

```
Out[57]: 1073741824
```

也可以通过API对内存池自定义或修改, 具体参见文档:

[Memory Management — CuPy 10.5.0 documentation](#)

CuPy和NumPy在某些行为上会有一些细微的不同, 包括:

- 浮点数转整数
- 随机方法
- 越界索引
- 重复索引处理
- 0维数组
- Matrix类型
- 数据类型
- UFUNC
- 随机种子
- NaN处理

具体可参见文档: [Differences between CuPy and NumPy — CuPy 10.5.0 documentation](#)

```
In [ ]:
```

CuPy可以和很多其他库结合使用, 比如NumPy, Numba, PyTorch等, 具体可参考文档:

[Interoperability — CuPy 10.5.0 documentation](#)

这部分内容涉及到了cuda编程，我们可能需要更灵活的控制，可以考虑使用PyCuda。

Sparse

文档: [Sparse — sparse 0.13.0+0.g0b7dfeb.dirty documentation](#)

`Sparse` 在NumPy和scipy.sparse上实现了任意维度的「稀疏数组」。

主要的数据结构参照遵循稀疏矩阵的Coordinate List (COO) 布局，并将其扩展到多个维度。

dmi1	dim2	dim3	...	data
0	0	0	.	10
0	0	3	.	13
0	2	2	.	9
3	1	4	.	21

除了存储，所有数组相关的操作（转置、reshape、切片、乘法等）都需要重新实现。

此外，本仓库还包括其他几个数据结构，比如Dictionary of Keys (DOK) 格式，它可以很好地推广到任意数量的维度。DOK适合编写和操作，但其他操作并不支持。常用的最佳实践是使用DOK编写一个数组，然后转为另一种格式执行其他操作。

也支持Compressed Sparse Row/Column (CSR/CSC) 格式。

创建

```
In [4]: import sparse as se
se.__version__
```

```
Out[4]: '0.13.0'
```

```
In [34]: rng = np.random.default_rng(42)
a = rng.random((100, 100, 100))
```

```
In [35]: # 构造一个稀疏矩阵，90%为0
a[a<0.9] = 0
```

```
In [36]: s = se.COO(a)
```

```
In [37]: s.nbytes
```

```
Out[37]: 3205760
```

```
In [38]: a.nbytes
```

```
Out[38]: 8000000
```

In [39]:

```
s
```

Out[39]:

Format	coo
---------------	-----

Data Type	float64
------------------	---------

Shape	(100, 100, 100)
--------------	-----------------

nnz	100180
------------	--------

Density	0.10018
----------------	---------

Read-only	True
------------------	------

Size	3.1M
-------------	------

Storage ratio	0.4
----------------------	-----

In [40]:

```
s.data
```

Out[40]:

```
array([0.97562235, 0.92676499, 0.97069802, ..., 0.98212211, 0.95084277,
       0.98694171])
```

In [41]:

```
s.coords
```

Out[41]:

```
array([[ 0,  0,  0, ..., 99, 99, 99],
       [ 0,  0,  0, ..., 99, 99, 99],
       [ 5, 11, 22, ..., 94, 95, 98]])
```

也可以直接通过坐标和值创建:

In [15]:

```
coords = [[0, 1, 2, 3, 4],
           [0, 1, 2, 3, 4]]
data = [10, 20, 30, 40, 50]
s = se.COO(coords, data, shape=(5, 5))
```

In [16]:

```
s
```

Out[16]:

Format	coo
---------------	-----

Data Type	int64
------------------	-------

Shape	(5, 5)
--------------	--------

nnz	5
------------	---

Density	0.2
----------------	-----

Read-only	True
------------------	------

Size	120
-------------	-----

Storage ratio	0.6
----------------------	-----

In [17]:

```
s.todense()
```

Out[17]:

```
array([[10,  0,  0,  0,  0],
       [ 0, 20,  0,  0,  0],
       [ 0,  0, 30,  0,  0],
       [ 0,  0,  0, 40,  0],
       [ 0,  0,  0,  0, 50]])
```

维度也可以比较随意:

```
In [23]: coords = [[0, 3, 2, 1], [4, 1, 2, 0]]
data = [1, 4, 2, 1]
s = se.COO(coords, data, shape=(6, 5))
s.todense()
```

```
Out[23]: array([[0, 0, 0, 0, 1],
               [1, 0, 0, 0, 0],
               [0, 0, 2, 0, 0],
               [0, 4, 0, 0, 0],
               [0, 0, 0, 0, 0],
               [0, 0, 0, 0, 0]])
```

```
In [24]: # 指定填充值
s = se.COO(coords, data, shape=(5,5), fill_value=-1)
s.todense()
```

```
Out[24]: array([[ -1,  -1,  -1,  -1,   1],
               [  1,  -1,  -1,  -1,  -1],
               [-1,  -1,   2,  -1,  -1],
               [-1,   4,  -1,  -1,  -1],
               [-1,  -1,  -1,  -1,  -1]])
```

从SciPy的稀疏矩阵、从NumPy的数组生成:

- `se.COO.from_scipy_sparse(x)`
- `se.COO.from_numpy(x)`

也支持随机:

```
In [28]: s = se.random((5, 5), density=0.1)
s
```

```
Out[28]: Format      coo
Data Type  float64
Shape      (5, 5)
nnz        2
Density    0.08
Read-only  True
Size       20
Storage ratio 0.1
```

```
In [29]: s.todense()
```

```
Out[29]: array([[0.         , 0.         , 0.1182219 , 0.         , 0.         ],
               [0.         , 0.         , 0.         , 0.         , 0.         ],
               [0.         , 0.         , 0.         , 0.         , 0.         ],
               [0.         , 0.         , 0.         , 0.         , 0.         ],
               [0.         , 0.         , 0.         , 0.         , 0.65295601]])
```

```
In [32]: s.data
```

```
Out[32]: array([0.1182219 , 0.65295601])
```

```
In [33]: s.coords
```

```
Out[33]: array([[0, 4],
               [2, 4]], dtype=uint8)
```

或者通过传入字典创建：

```
In [42]: d = {(0, 0, 0): 1, (1, 2, 3): 2, (1, 1, 0): 3}
s = se.COO(d)
```

```
In [44]: s.shape
```

```
Out[44]: (2, 3, 4)
```

```
In [45]: s
```

```
Out[45]: 

|                      |           |
|----------------------|-----------|
| <b>Format</b>        | coo       |
| <b>Data Type</b>     | int64     |
| <b>Shape</b>         | (2, 3, 4) |
| <b>nnz</b>           | 3         |
| <b>Density</b>       | 0.125     |
| <b>Read-only</b>     | True      |
| <b>Size</b>          | 96        |
| <b>Storage ratio</b> | 0.5       |


```

也可以通过数组：

```
In [48]: L = [((0, 0), 1),
              ((1, 1), 2),
              ((0, 0), 3)]
```

```
In [50]: s = se.COO(L)
s.todense()
```

```
Out[50]: array([[4, 0],
               [0, 2]])
```

或者从DOK转换过来：

```
In [46]: s1 = se.DOK((5, 5))
s1
```

```
Out[46]: Format      dok
Data Type float64
Shape      (5, 5)
nnz        0
Density     0.0
Read-only   False
Size        0
Storage ratio 0.0
```

```
In [51]: s1[1:3, 1:3] = [[4,5],[6,7]]
```

```
In [52]: s1.todense()
```

```
Out[52]: array([[0., 0., 0., 0., 0.],
               [0., 4., 5., 0., 0.],
               [0., 6., 7., 0., 0.],
               [0., 0., 0., 0., 0.],
               [0., 0., 0., 0., 0.]])
```

```
In [59]: s1.coords
```

```
-----
AttributeError                                Traceback (most recent call last)
Input In [59], in <cell line: 1>()
----> 1 s1.coords
AttributeError: 'DOK' object has no attribute 'coords'
```

```
In [56]: s1.data
```

```
Out[56]: {(1, 1): 4.0, (1, 2): 5.0, (2, 1): 6.0, (2, 2): 7.0}
```

```
In [53]: s2 = s1.asformat("coo")
```

```
In [54]: s2.todense()
```

```
Out[54]: array([[0., 0., 0., 0., 0.],
               [0., 4., 5., 0., 0.],
               [0., 6., 7., 0., 0.],
               [0., 0., 0., 0., 0.],
               [0., 0., 0., 0., 0.]])
```

```
In [58]: s2.coords
```

```
Out[58]: array([[1, 1, 2, 2],
               [1, 2, 1, 2]])
```

```
In [57]: s2.data
```

```
Out[57]: array([4., 5., 6., 7.])
```

```
In [61]: # 这样也可以转换
s3 = se.COO(s1)
```



```
s3.todense()
```

```
Out[61]: array([[0., 0., 0., 0., 0.],
               [0., 4., 5., 0., 0.],
               [0., 6., 7., 0., 0.],
               [0., 0., 0., 0., 0.],
               [0., 0., 0., 0., 0.]])
```

转换

C00 对象可以转换成其他格式，包括：

- `C00.todense`：转成NumPy数组
- `C00.mayhbe_densify`：基于某些条件转为NumPy数组
- `C00.to_scipy_sparse`：如果数组是二维，则转成 `spicy.sparse.coo_matrix`
- `C00.tocsr`：如果数组是二维，转成 `scipy.sparse.csr_matrix`
- `C00.tocsc`：如果数组是二维，转成 `scipy.sparse.csc_matrix`

着重说一下第二个API，它接受两个参数：`max_size`（输出的最大元素数，默认1000）和`min_density`（输出的最小密度，默认0.25），当一个稀疏数组两个条件都不满足时，会抛出异常。

```
In [128... x = np.zeros((5, 5), dtype=np.uint8)
x[2, :] = 1
s = se.C00.from_numpy(x)
```

```
In [129... x
```

```
Out[129... array([[0, 0, 0, 0, 0],
        [0, 0, 0, 0, 0],
        [1, 1, 1, 1, 1],
        [0, 0, 0, 0, 0],
        [0, 0, 0, 0, 0]], dtype=uint8)
```

```
In [146... # 25满足，0.9不满足
s.maybe_densify(max_size=25, min_density=0.21)
```

```
Out[146... array([[0, 0, 0, 0, 0],
        [0, 0, 0, 0, 0],
        [1, 1, 1, 1, 1],
        [0, 0, 0, 0, 0],
        [0, 0, 0, 0, 0]], dtype=uint8)
```

```
In [147... # 24不满足，0.1满足
s.maybe_densify(max_size=24, min_density=0.1)
```

```
Out[147... array([[0, 0, 0, 0, 0],
        [0, 0, 0, 0, 0],
        [1, 1, 1, 1, 1],
        [0, 0, 0, 0, 0],
        [0, 0, 0, 0, 0]], dtype=uint8)
```

```
In [148... # 都满足
s.maybe_densify(max_size=25, min_density=0.1)
```

```
Out[148...] array([[0, 0, 0, 0, 0],
          [0, 0, 0, 0, 0],
          [1, 1, 1, 1, 1],
          [0, 0, 0, 0, 0],
          [0, 0, 0, 0, 0]], dtype=uint8)
```

```
In [149...] # 都不满足
s.maybe_densify(max_size=24, min_density=0.21)
```

```
-----
ValueError                                Traceback (most recent call last)
Input In [149], in <cell line: 2>()
      1 # 都不满足
----> 2 s.maybe_densify(max_size=24, min_density=0.21)

File ~/home/env/anaconda3/envs/tf29/lib/python3.8/site-packages/sparse/_coo/core.p
y:1379, in COO.maybe_densify(self, max_size, min_density)
    1377     return self.todense()
    1378 else:
-> 1379     raise ValueError(
    1380         "Operation would require converting " "large sparse array to dens
e"
    1381     )

ValueError: Operation would require converting large sparse array to dense
```

计算

```
In [162...] s = se.random((3, 3, 3), density=0.1)
s.todense()
```

```
Out[162...] array([[[0.        , 0.        , 0.        ],
          [0.        , 0.        , 0.        ],
          [0.        , 0.        , 0.        ]],

          [[0.        , 0.        , 0.        ],
          [0.82191339, 0.        , 0.        ],
          [0.        , 0.        , 0.        ]],

          [[0.        , 0.        , 0.        ],
          [0.        , 0.        , 0.        ],
          [0.        , 0.312388 , 0.        ]])
```

```
In [163...] y = np.sin(s) + s.T * 1
y
```

Out[163...	Format	COO
	Data Type	float64
	Shape	(3, 3, 3)
	nnz	4
	Density	0.14814814814814814
	Read-only	True
	Size	44
	Storage ratio	0.2

In [164... `y.todense()`

Out[164... `array([[[0. , 0. , 0.],`
 `[0. , 0.82191339, 0.],`
 `[0. , 0. , 0.]],`

 `[[0. , 0. , 0.],`
 `[0.73244985, 0. , 0.],`
 `[0. , 0. , 0.312388]],`

 `[[0. , 0. , 0.],`
 `[0. , 0. , 0.],`
 `[0. , 0.30733194, 0.]]])`

更多内容可参考：

[Operations on COO and GCXS arrays — sparse 0.13.0+0.g0b7dfef.dirty documentation](#)

Dask

文档：[Dask — Dask documentation](#)

Dask主要用于并行计算。包括两部分：

- 针对计算优化的动态任务调度，类似于Airflow, Luigi, Celery或Make，但针对交互式计算工作负载进行了优化。
- 大数据集合，如并行数组，DataFrame，列表等，将NumPy, Pandas或Python迭代器等常见接口扩展到大于内存或分布式环境。这些并行集合在动态任务调度程序上运行。

整体架构如下：

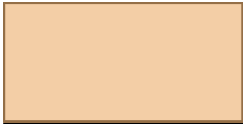
Dask可以处理很多数据集合，我们这里以数组（Array）为例。关于Dask，我们在第一节《数组对象》也略有提及。

In [15]: `import dask.array as da`

```
In [4]: a = np.array([2,3], like=da.array([1,1]))
a
```

Out[4]:

	Array	Chunk
Bytes	16 B	16 B
Shape	(2,)	(2,)
Count	1 Tasks	1 Chunks
Type	int64	numpy.ndarray



```
In [5]: a.compute()
```

Out[5]: array([2, 3])

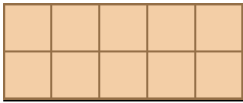
创建

```
In [6]: data = np.arange(100_000).reshape(200, 500)
```

```
In [7]: a = da.from_array(data, chunks=(100, 100))
a
```

Out[7]:


	Array	Chunk
Bytes	781.25 kiB	78.12 kiB
Shape	(200, 500)	(100, 100)
Count	10 Tasks	10 Chunks
Type	int64	numpy.ndarray



```
In [8]: # 自动 chunk
b = da.from_array(data)
b
```

Out[8]:

	Array	Chunk
Bytes	781.25 kiB	781.25 kiB
Shape	(200, 500)	(200, 500)
Count	1 Tasks	1 Chunks
Type	int64	numpy.ndarray




索引

```
In [9]: a[0]
```

Out[9]:


	Array	Chunk
Bytes	3.91 kiB	800 B
Shape	(500,)	(100,)
Count	15 Tasks	5 Chunks
Type	int64	numpy.ndarray



In [10]: `a[:50, 200]`

Out[10]:

	Array	Chunk
Bytes	400 B	400 B
Shape	(50,)	(50,)
Count	11 Tasks	1 Chunks
Type	int64	numpy.ndarray



计算

Dask是惰性估计的，会生成一个用于计算的Dask任务图，请求结果时才会计算。计算和NumPy类似，也可以结合NumPy。

In [11]: `a.compute()`

Out[11]: `array([[0, 1, 2, ..., 497, 498, 499],
[500, 501, 502, ..., 997, 998, 999],
[1000, 1001, 1002, ..., 1497, 1498, 1499],
...,
[98500, 98501, 98502, ..., 98997, 98998, 98999],
[99000, 99001, 99002, ..., 99497, 99498, 99499],
[99500, 99501, 99502, ..., 99997, 99998, 99999]])`

In [12]: `a.mean()`

Out[12]:

	Array	Chunk
Bytes	8 B	8.0 B
Shape	()	()
Count	26 Tasks	1 Chunks
Type	float64	numpy.ndarray


In [13]: `a.mean().compute()`

Out[13]: `49999.5`

In [14]: `np.sin(a)`

Out[14]:

	Array	Chunk
Bytes	781.25 kiB	78.12 kiB
Shape	(200, 500)	(100, 100)
Count	20 Tasks	10 Chunks
Type	float64	numpy.ndarray



In [15]: `_.compute()`

Out[15]: array([[0. , 0.84147098, 0.90929743, ..., 0.58781939,
 0.99834363, 0.49099533],
 [-0.46777181, -0.9964717 , -0.60902011, ..., -0.89796748,
 -0.85547315, -0.02646075],
 [0.82687954, 0.9199906 , 0.16726654, ..., 0.99951642,
 0.51387502, -0.4442207],
 ...,
 [-0.99720859, -0.47596473, 0.48287891, ..., -0.76284376,
 0.13191447, 0.90539115],
 [0.84645538, 0.00929244, -0.83641393, ..., 0.37178568,
 -0.5802765 , -0.99883514],
 [-0.49906936, 0.45953849, 0.99564877, ..., 0.10563876,
 0.89383946, 0.86024828]])


任务图

In [20]: `c = a.max(axis=1)[: -1] + 10`

In [21]: `c`

Out[21]:

	Array	Chunk
Bytes	1.56 kiB	800 B
Shape	(200,)	(100,)
Count	30 Tasks	2 Chunks
Type	int64	numpy.ndarray



In [23]: `c.dask`

Out[23]:



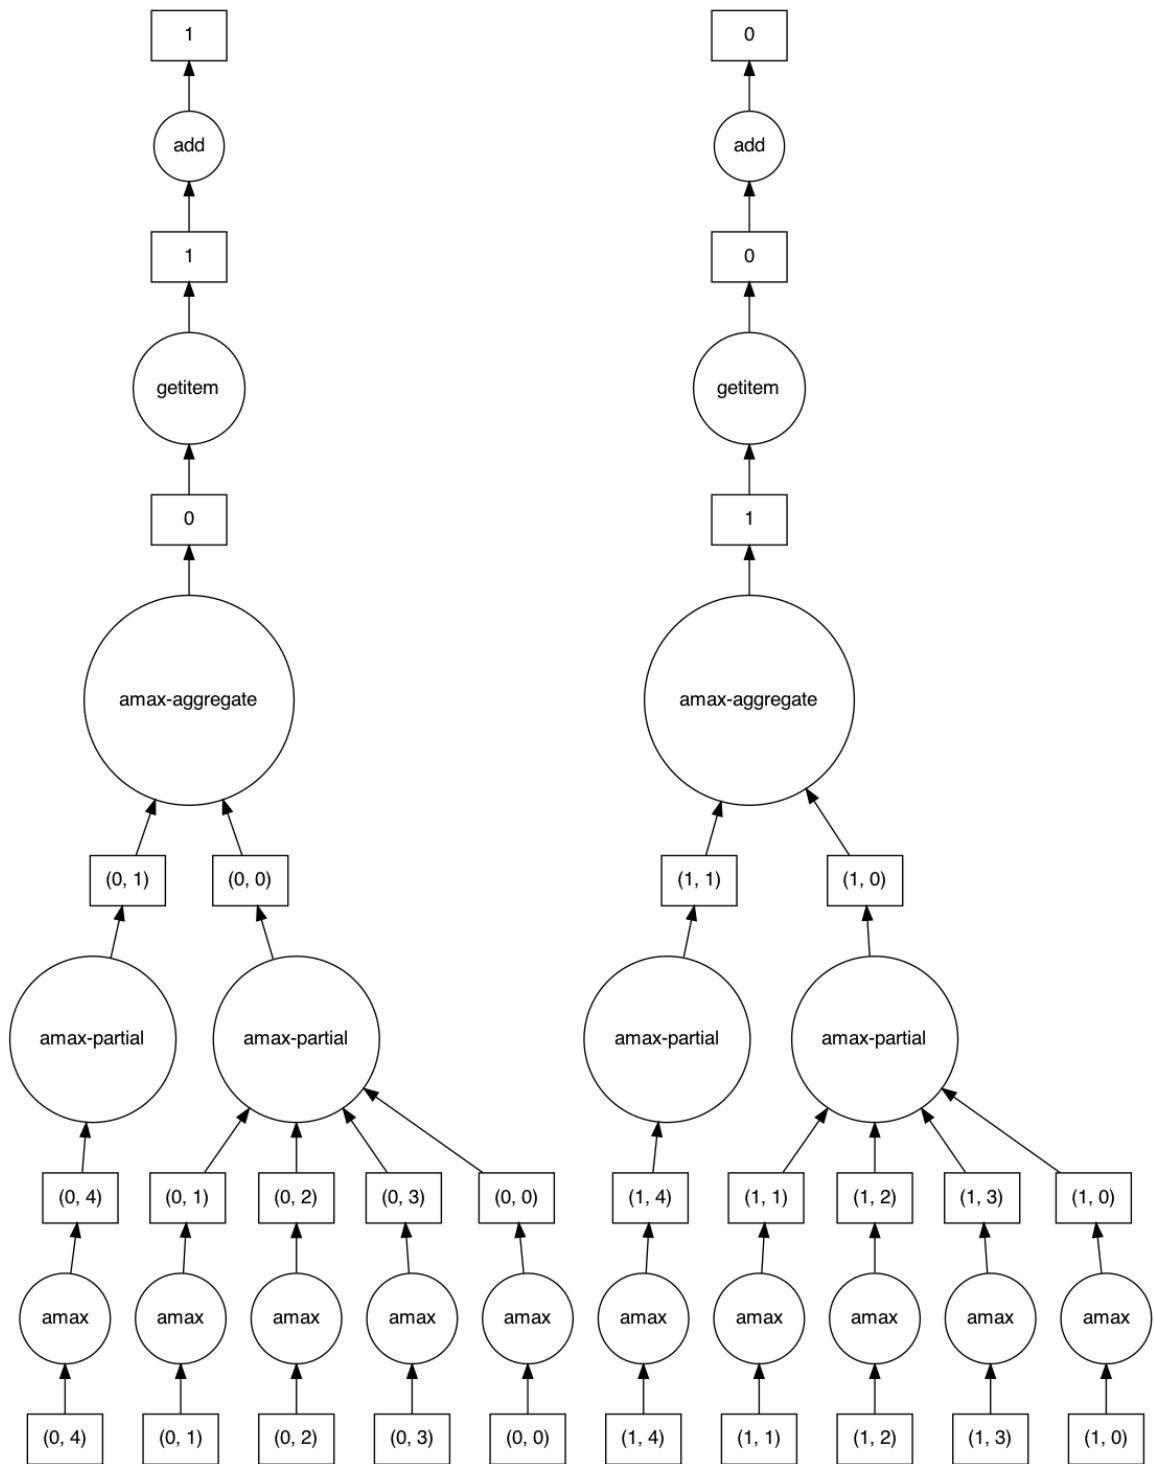
HighLevelGraph

HighLevelGraph with 6 layers and 30 keys from all layers.

- ▶ Layer1: array
- ▶ Layer2: amax
- ▶ Layer3: amax-partial
- ▶ Layer4: amax-aggregate
- ▶ Layer5: getitem
- ▶ Layer6: add

In [24]: `c.visualize()`

Out[24]:



底层API

使用 `delayed` 装饰器将函数调用包装到一个延迟构造的任务图中：

```
In [1]: import dask
```

```
In [6]: @dask.delayed
def inc(x):
    return x+1
@dask.delayed
def add(x, y):
    return x+y
```



```
In [10]: a = inc(1)
         b = inc(2)
         c = add(a, b)
```

```
In [12]: a, b, c
```

```
Out[12]: (Delayed('inc-171697d8-9c8e-42a6-b8f2-3a1c5ff36faa'),
          Delayed('inc-5a5d5200-1e59-4d9d-b0e5-049cad501cf4'),
          Delayed('add-3d37a606-f83d-401e-bacb-c16f9bec7c9a'))
```

```
In [13]: c = c.compute()
         c
```

```
Out[13]: 5
```

调度

在生成任务图后，调度程序默认使用计算机线程池运行计算。

线程调度使用本地的 `concurrent.futures.ThreadPoolExecutor` 执行计算，是Dask Array，Dask DataFrame和Dask Delayed的默认选择。

由于Python的全局解释器锁 (GIL)，此调度程序仅在计算由非Python代码主导时提供并行性。主要是在NumPy数组、Pandas DataFrames中的数字数据上操作或使用任何生态系统中其他基于C/C++/Cython的项目。

```
In [3]: # 对scheduler进行配置
        dask.config.set(scheduler="threads")
```

```
Out[3]: <dask.config.set at 0x104bfdbe0>
```

```
In [5]: dask.config.get("scheduler")
```

```
Out[5]: 'threads'
```

进程调度使用本地的 `concurrent.futures.ProcessPoolExecutor` 执行计算，是Dask Bag的默认选择。

每个任务及其所有依赖项都被传送到本地进程执行，然后它们的结果被传送回主进程。可以绕过Python的GIL问题。但是将数据移动到进程可能会导致性能下降，尤其是在进程间传输大量数据时。当不涉及任务间数据传输，输出输入都很小时是一个很好的选择。

```
In [6]: dask.config.set(scheduler='processes')
```

```
Out[6]: <dask.config.set at 0x1046037c0>
```

```
In [7]: dask.config.get("scheduler")
```

```
Out[7]: 'processes'
```

单线程同步调度器在一个本地线程中执行所有的计算，没有并行。一般用于调试或分析。比如Jupyter Notebook的魔法方法 `%debug`，`%pdb`，`%prun` 等在使用并行Dask调度时将无法正常工作。

```
In [8]: dask.config.set(scheduler='synchronous')
```

```
Out[8]: <dask.config.set at 0x104bfd9d0>
```

```
In [9]: dask.config.get("scheduler")
```

```
Out[9]: 'synchronous'
```

Dask支持使用分布式调度器进行更多控制，它可以在单台或多台机器上工作，可将其视为「高级调度器」。这也是目前比较推荐的方式。


之所以在单台机器上也推荐使用，原因在于：

- 提供异步API访问，尤其是Futures
- 提供诊断仪表板，可以提供有关性能和进度的意见
- 以更复杂的方式处理数据的局部性，在需要多个进程的工作负载上比多处理器调度器更有效

```
In [10]: from dask.distributed import Client
```

```
In [11]: client = Client()
```

```
In [12]: client
```

```
Out[12]:  Client
Client-81d68614-e4ad-11ec-9f7c-acde48001122
Connection method: Cluster object    Cluster type: distributed.LocalCluster
Dashboard:
http://127.0.0.1:8787/status
```

► Cluster Info

```
In [13]: dask.config.get("scheduler")
```

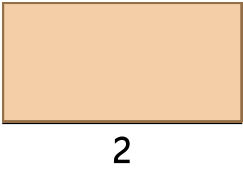
```
Out[13]: 'dask.distributed'
```

请注意上面 Dashboard 的地址，可以直接访问。

当然了，除了上面的这种全局配置方法，也可以使用上下文管理器，或执行 `compute` 时通过参数传入。

```
In [18]: x = da.array([1,2])
x
```

Out[18]:

	Array	Chunk		
Bytes	16 B	16 B		
Shape	(2,)	(2,)		
Count	1 Tasks	1 Chunks		
Type	int64	numpy.ndarray		

```
In [21]: with dask.config.set(scheduler="threads"):
         xo = x.compute()
```

In [22]: xo

Out[22]: array([1, 2])

```
In [23]: dask.config.get("scheduler")
```

Out[23]: 'dask.distributed'

```
In [25]: x.compute(scheduler="threads")
```

Out[25]: array([1, 2])

更多关于分布式调度器的使用可参阅文档：

[Deploy Dask Clusters — Dask documentation](#)

性能对比

```
In [7]: %%time
         rng = np.random.default_rng(42)
         x = rng.normal(10, 0.1, size=(20000, 20000))
         y = x.mean(axis=0)[:100]
```

CPU times: user 8.14 s, sys: 894 ms, total: 9.03 s
Wall time: 9.48 s

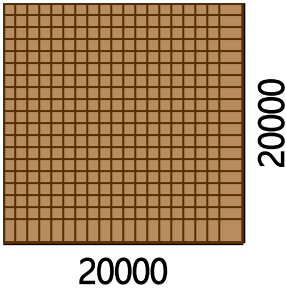
```
In [8]: %%time
         x = da.random.normal(
             10, 0.1, size=(20000, 20000), chunks=(1000, 1000)
         )
         y = x.mean(axis=0)[:100]
         o = y.compute()
```

CPU times: user 26.1 s, sys: 596 ms, total: 26.7 s
Wall time: 8.75 s

```
In [9]: x
```

Out[9]:

	Array	Chunk
Bytes	2.98 GiB	7.63 MiB
Shape	(20000, 20000)	(1000, 1000)
Count	400 Tasks	400 Chunks
Type	float64	numpy.ndarray



Dask完成的更快，但使用了更多的总CPU时间。

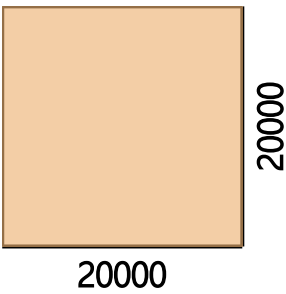
```
In [43]: %%time
x = da.random.normal(
    10, 0.1, size=(20000, 20000), chunks=(20000, 20000)
)
y = x.mean(axis=0)[:100]
o = y.compute()
```

CPU times: user 39.3 s, sys: 8.38 s, total: 47.7 s
Wall time: 59.3 s

```
In [44]: x
```

Out[44]:

	Array	Chunk
Bytes	2.98 GiB	2.98 GiB
Shape	(20000, 20000)	(20000, 20000)
Count	1 Tasks	1 Chunks
Type	float64	numpy.ndarray




```
In [41]: %%time
x = da.random.normal(
    10, 0.1, size=(20000, 20000), chunks=(25, 25)
)
y = x.mean(axis=0)[:100]
o = y.compute()
```

CPU times: user 3min 30s, sys: 36.4 s, total: 4min 7s
Wall time: 4min 16s

```
In [39]: x
```

Out[39]:

	Array	Chunk
Bytes	2.98 GiB	4.88 kiB
Shape	(20000, 20000)	(25, 25)
Count	640000 Tasks	640000 Chunks
Type	float64	numpy.ndarray



Xarray

文档: [xarray: N-D labeled arrays and datasets in Python](#)

xarray在原始的类NumPy数组之上以维度、坐标和属性的形式引入标签，从而提供更直观、更简洁且不易出错的开发体验。

与之相关的工具包可参见：

Installation

xarray有两个核心的数据结构，它们构建在NumPy和Pandas之上，并进行了扩展，都是多维的：

- DataArray：有标签的N维数组
- Dataset：多维内存数组数据库

```
In [2]: import xarray as xr
xr.__version__
```

Out[2]: '2022.3.0'

创建


```
In [106... rng = np.random.default_rng(42)
a = rng.normal(size=(2,3))
a
```

```
Out[106... array([[ 0.30471708, -1.03998411,  0.7504512 ],
        [ 0.94056472, -1.95103519, -1.30217951]])
```

```
In [107... data = xr.DataArray(a, dims=("x", "y"), coords={"x": [10, 20]})
```

```
In [108... data
```

Out[108...] xarray.DataArray (x: 2, y: 3)

 array([[0.30471708, -1.03998411, 0.7504512],
 [0.94056472, -1.95103519, -1.30217951]])

▼ Coordinates:


x (x) int64 10 20



► Attributes: (0)

In [109...] data.x

Out[109...] xarray.DataArray 'x' (x: 2)

 array([10, 20])

▼ Coordinates:

x (x) int64 10 20



► Attributes: (0)

In [110...] data.y

Out[110...] xarray.DataArray 'y' (y: 3)

 array([0, 1, 2])

► Coordinates: (0)

► Attributes: (0)

In [111...] data.values

Out[111...] array([[0.30471708, -1.03998411, 0.7504512],
 [0.94056472, -1.95103519, -1.30217951]])

In [112...] data.dims

Out[112...] ('x', 'y')

In [113...] data.coords

Out[113...] Coordinates:
* x (x) int64 10 20

In [114...] data.attrs

Out[114...] {}

In [115...] data.x.values

Out[115...] array([10, 20])


In [116...] data.y.values

Out[116... array([0, 1, 2])

索引

In [60]: data[0,:]

Out[60]: xarray.DataArray (y: 3)

 array([0.30471708, -1.03998411, 0.7504512])

▼ Coordinates:


x () int64 10



► Attributes: (0)

In [63]: data.loc[10]

Out[63]: xarray.DataArray (y: 3)

 array([0.30471708, -1.03998411, 0.7504512])

▼ Coordinates:

x () int64 10



► Attributes: (0)

In [68]: data.loc[0]

```

-----
KeyError                                Traceback (most recent call last)
/usr/local/lib/python3.8/site-packages/pandas/core/indexes/base.py in get_loc(self, key, method, tolerance)
    3620         try:
-> 3621             return self._engine.get_loc(casted_key)
    3622         except KeyError as err:

/usr/local/lib/python3.8/site-packages/pandas/_libs/index.pyx in pandas._libs.index.IndexEngine.get_loc()

/usr/local/lib/python3.8/site-packages/pandas/_libs/index.pyx in pandas._libs.index.IndexEngine.get_loc()

pandas/_libs/hashtable_class_helper.pxi in pandas._libs.hashtable.Int64HashTable.get_item()

pandas/_libs/hashtable_class_helper.pxi in pandas._libs.hashtable.Int64HashTable.get_item()

KeyError: 0

```

The above exception was the direct cause of the following exception:

```

KeyError                                Traceback (most recent call last)
<ipython-input-68-0dc99c936dd2> in <module>
----> 1 data.loc[0]

/usr/local/lib/python3.8/site-packages/xarray/core/dataarray.py in __getitem__(self, key)
    197         labels = indexing.expanded_indexer(key, self.data_array.ndim)
    198         key = dict(zip(self.data_array.dims, labels))
--> 199         return self.data_array.sel(key)
    200
    201     def __setitem__(self, key, value) -> None:

/usr/local/lib/python3.8/site-packages/xarray/core/dataarray.py in sel(self, indexers, method, tolerance, drop, **indexers_kwargs)
    1327         Dimensions without coordinates: points
    1328         """
-> 1329         ds = self._to_temp_dataset().sel(
    1330             indexers=indexers,
    1331             drop=drop,

/usr/local/lib/python3.8/site-packages/xarray/core/dataset.py in sel(self, indexers, method, tolerance, drop, **indexers_kwargs)
    2499         """
    2500         indexers = either_dict_or_kwargs(indexers, indexers_kwargs, "sel")
-> 2501         pos_indexers, new_indexes = remap_label_indexers(
    2502             self, indexers=indexers, method=method, tolerance=tolerance
    2503         )

/usr/local/lib/python3.8/site-packages/xarray/core/coordinates.py in remap_label_indexers(obj, indexers, method, tolerance, **indexers_kwargs)
    419     }
    420
--> 421     pos_indexers, new_indexes = indexing.remap_label_indexers(
    422         obj, v_indexers, method=method, tolerance=tolerance
    423     )

```



```

/usr/local/lib/python3.8/site-packages/xarray/core/indexing.py in remap_label_indexers(data_obj, indexers, method, tolerance)
    119     for dim, index in indexes.items():
    120         labels = grouped_indexers[dim]
--> 121         idxr, new_idx = index.query(labels, method=method, tolerance=tolerance)
    122         pos_indexers[dim] = idxr
    123         if new_idx is not None:

/usr/local/lib/python3.8/site-packages/xarray/core/indexes.py in query(self, labels, method, tolerance)
    239             )
    240             else:
--> 241                 indexer = self.index.get_loc(label_value)
    242             elif label.dtype.kind == "b":
    243                 indexer = label

/usr/local/lib/python3.8/site-packages/pandas/core/indexes/base.py in get_loc(self, key, method, tolerance)
    3621         return self._engine.get_loc(casted_key)
    3622     except KeyError as err:
-> 3623         raise KeyError(key) from err
    3624     except TypeError:
    3625         # If we have a listlike key, _check_indexing_error will raise
KeyError: 0

```

```
In [70]: # integer select
data.isel(x=0)
```

```
Out[70]: xarray.DataArray (y: 3)
```

```
array([ 0.30471708, -1.03998411,  0.7504512 ])
```

▼ Coordinates:

```
x          () int64 10
```

► Attributes: (0)

```
In [71]: data.isel(y=0)
```

```
Out[71]: xarray.DataArray (x: 2)
```

```
array([0.30471708, 0.94056472])
```


▼ Coordinates:

```
x          (x) int64 10 20
```

► Attributes: (0)

```
In [73]: # 直接select
data.sel(x=10)
```

Out[73]: xarray.DataArray (y: 3)

 array([0.30471708, -1.03998411, 0.7504512])

▼ Coordinates:

x () int64 10

► Attributes: (0)

更多介绍可参阅:

[Indexing and selecting data](#)

属性


在设置DataArray时, 设置元数据属性通常是个不错的实践。常见的属性包括

`long_name`, `units` 等。

```
In [117... data.attrs["long_name"] = "random velocity"
data.attrs["units"] = "metres/sec"
data.attrs["description"] = "A random variable created as an example."
```

In [118... data

Out[118... xarray.DataArray (x: 2, y: 3)

 array([[0.30471708, -1.03998411, 0.7504512],
 [0.94056472, -1.95103519, -1.30217951]])

▼ Coordinates:

x (x) int64 10 20

▼ Attributes:

long_name : random velocity
units : metres/sec
description : A random variable created as an example.


In [119... data.attrs

Out[119... {'long_name': 'random velocity',
 'units': 'metres/sec',
 'description': 'A random variable created as an example.'}

```
In [120... # 给坐标设置属性
data.x.attrs["units"] = "x units"
```

In [121... data.x

Out[121... xarray.DataArray 'x' (x: 2)

 array([10, 20])

▼ Coordinates:

x (x) int64 10 20



▼ Attributes:


units : x units



计算

In [94]: data + 10

Out[94]: xarray.DataArray (x: 2, y: 3)

 array([[10.30471708, 8.96001589, 10.7504512],
 [10.94056472, 8.04896481, 8.69782049]])

▼ Coordinates:


x (x) int64 10 20



► Attributes: (0)

In [95]: np.sin(data)

Out[95]: xarray.DataArray (x: 2, y: 3)

 array([[0.3000233 , -0.86239618, 0.68196883],
 [0.80789103, -0.92857601, -0.96413891]])

▼ Coordinates:

x (x) int64 10 20



▼ Attributes:

long_name : random velocity

units : metres/sec

description : A random variable created as an example.



In [96]: data.T

Out[96]: xarray.DataArray (y: 3, x: 2)

```
array([[ 0.30471708,  0.94056472],
       [-1.03998411, -1.95103519],
       [ 0.7504512 , -1.30217951]])
```

▼ Coordinates:

x	(x)	int64	10	20
---	-----	-------	----	----

▼ Attributes:

long_name :	random velocity
units :	metres/sec
description :	A random variable created as an example.

In [97]: data.sum()

Out[97]: xarray.DataArray

```
array(-2.29746581)
```

► Coordinates: (0)

► Attributes: (0)

In [98]: data.mean(dim="x")

Out[98]: xarray.DataArray (y: 3)

```
array([ 0.6226409 , -1.49550965, -0.27586416])
```

► Coordinates: (0)

► Attributes: (0)

与NumPy对比一下:

In [100... data.mean(axis=0)

Out[100... xarray.DataArray (y: 3)

```
array([ 0.6226409 , -1.49550965, -0.27586416])
```

► Coordinates: (0)

► Attributes: (0)

In [99]: data.mean(dim="y")

Out[99]: xarray.DataArray (x: 2)

 array([0.00506139, -0.77088333])

▼ Coordinates:

x (x) int64 10 20



► Attributes: (0)

基于维度名称的广播（不需要插入虚拟尺寸进行对齐）：

In [15]: `[data.coords["y"]]`

Out[15]: `[<xarray.DataArray 'y' (y: 3)>
array([0, 1, 2])
Dimensions without coordinates: y]`

In [28]: `rng = np.random.default_rng(42)
a = xr.DataArray(rng.integers(0, 10, 3), [data.coords["y"]])
b = xr.DataArray(rng.integers(0, 10, 4), dims="z")`

In [29]: `a`

Out[29]: xarray.DataArray (y: 3)

 array([0, 7, 6])

▼ Coordinates:


y (y) int64 0 1 2



► Attributes: (0)

In [30]: `b`

Out[30]: xarray.DataArray (z: 4)

 array([4, 4, 8, 0])

► Coordinates: (0)

► Attributes: (0)

In [31]: `a + b`

Out[31]: xarray.DataArray (y: 3, z: 4)

```
array([[ 4,  4,  8,  0],
       [11, 11, 15,  7],
       [10, 10, 14,  6]])
```

▼ Coordinates:

y (y) int64 0 1 2

► Attributes: (0)

In [32]: a.values + b.values

```
-----
ValueError                                Traceback (most recent call last)
<ipython-input-32-db9ccfe54ec4> in <module>
----> 1 a.values + b.values

ValueError: operands could not be broadcast together with shapes (3,) (4,)
```

意味着大多数情况下不需要担心维度顺序:

In [13]: data.shape

Out[13]: (2, 3)

In [12]: data - data.T

Out[12]: xarray.DataArray (x: 2, y: 3)

```
array([[0., 0., 0.],
       [0., 0., 0.]])
```

▼ Coordinates:

x (x) int64 10 20

► Attributes: (0)


更多可参阅:

[Computation](#)

GroupBy

```
In [82]: rng = np.random.default_rng(42)
a = xr.DataArray(rng.integers(1, 10, (2, 3)), dims=("x", "y"))
a
```

Out[82]: xarray.DataArray (x: 2, y: 3)


 array([[1, 7, 6],
[4, 4, 8]])

► Coordinates: (0)

► Attributes: (0)

```
In [83]: labels = xr.DataArray(["E", "F", "E"], dims="y", name="labels")  
labels
```

Out[83]: xarray.DataArray 'labels' (y: 3)


 array(['E', 'F', 'E'], dtype='<U1')

► Coordinates: (0)

► Attributes: (0)

```
In [85]: a.groupby(labels).mean("y")
```

Out[85]: xarray.DataArray (x: 2, labels: 2)

 array([[3.5, 7.],
[6. , 4.]])

▼ Coordinates:


labels (labels) object 'E' 'F'



► Attributes: (0)

```
In [86]: a.groupby(labels).mean("x")
```

Out[86]: xarray.DataArray (y: 3)


 array([2.5, 5.5, 7.])

► Coordinates: (0)

► Attributes: (0)

```
In [87]: # 1 4 6 8 一组  
# 7 4 一组  
a.groupby(labels).map(lambda x: x - x.min())
```

Out[87]: xarray.DataArray (x: 2, y: 3)

 array([[0, 3, 5],
[3, 0, 7]])

► Coordinates: (0)

► Attributes: (0)


可视化

```
In [125... rng = np.random.default_rng(42)
a = xr.DataArray(
    rng.integers(1, 10, (2, 3)),
    dims=("x", "y"),
    coords={"x": [10, 20], "y": [10, 20, 30]},
    attrs={
        "long_name": "random integers",
        "units": "null",
        "description": "demo random"
    }
)
```





```
In [126... a.x.attrs["units"] = "x 10-20"
```

```
In [127... a
```

Out[127... xarray.DataArray (x: 2, y: 3)

 array([[1, 7, 6],
[4, 4, 8]])

▼ Coordinates:

x	(x) int64 10 20	 
y	(y) int64 10 20 30	 


▼ Attributes:

long_name : random integers
units : null
description : demo random



```
In [128... a.plot()
```

Out[128... <matplotlib.collections.QuadMesh at 0x11b0604f0>



小结

In []:

参考

- [Beyond Numpy Arrays in Python](#)

In []: